

How to Use Survival Analytics to Predict Employee Turnover

(Or, Why You Shouldn't Use Logistic Regression to Predict Attrition)

Talent Analytics, Corp.

Pasha Roberts, Chief Scientist

May 16, 2017

Demo Code: https://github.com/talentanalytics/class_survival_101/



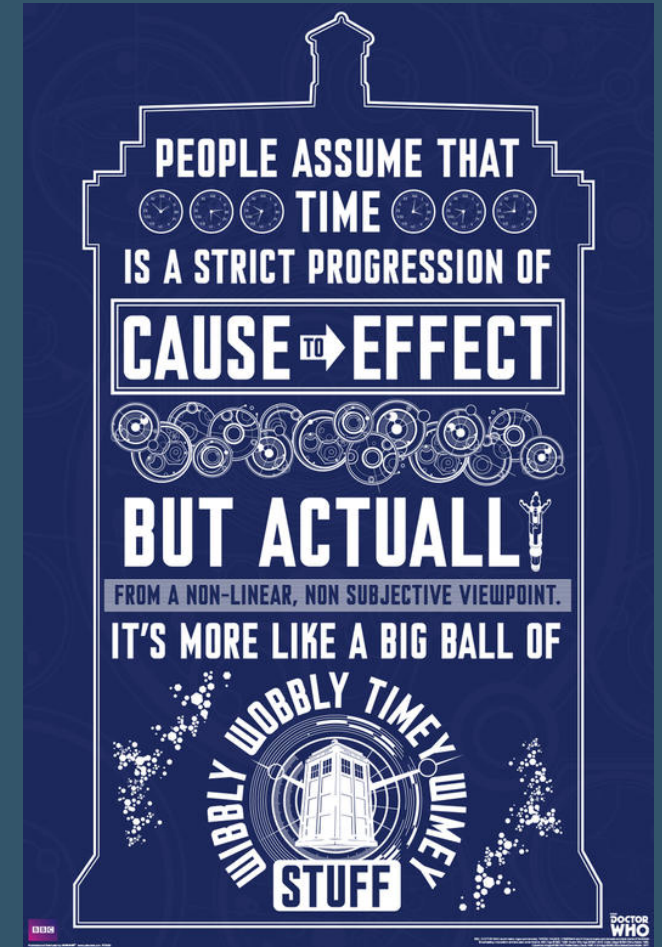
Who is Talent Analytics?

Predictive Modeling Platform – Advisor™

- ▶ We predict employee attrition and performance pre-hire.
- ▶ Much like credit risk modeling:
 - ▷ Predict likelihood to pay / default on mortgage, **before** extending credit
 - ▷ Predict likelihood to perform / leave role or company early, **before** extending job offer
- ▶ PAAS (Prediction As A Service)
- ▶ Seamless deployment of predictions into talent acquisition process




Turnover and Tenure: A Timey-Wimey Relationship



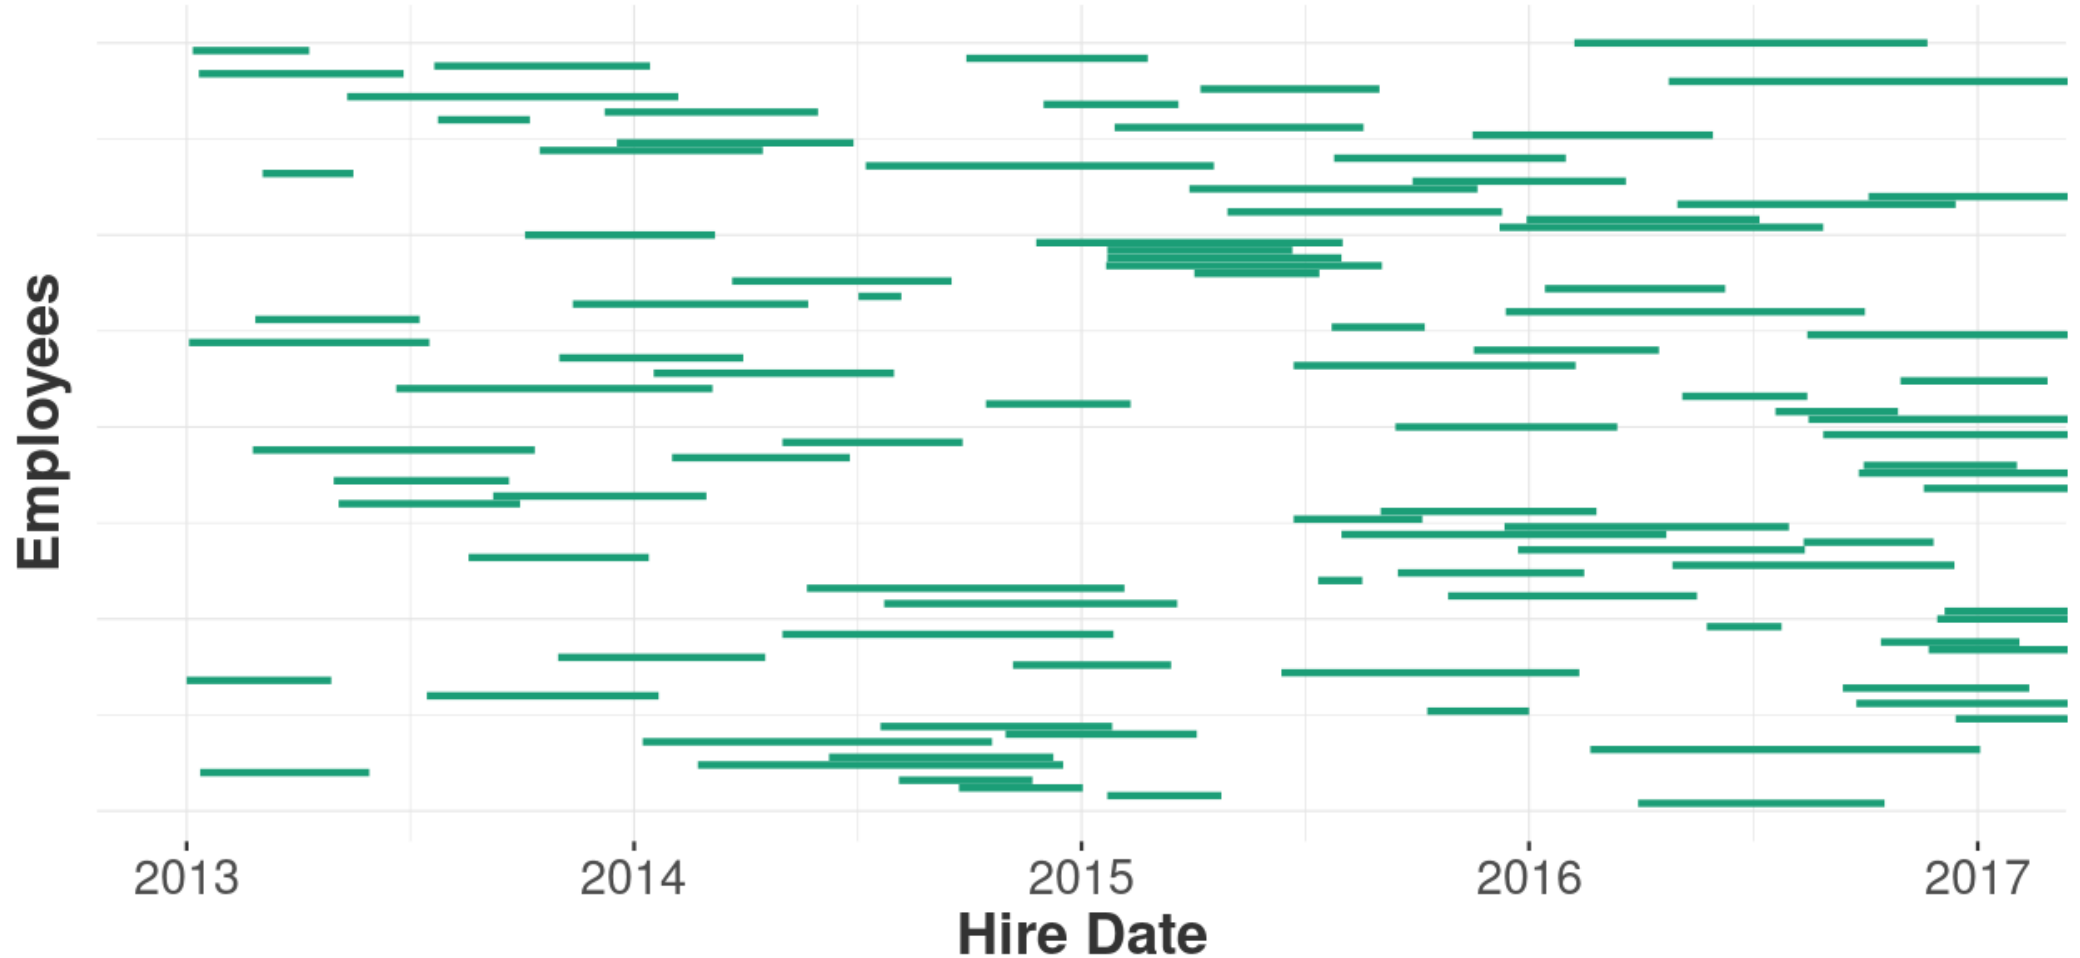
Two Different (But Related) Things

- ▶ **Turnover:** Percent of employees that terminate within a period of time
- ▶ **Tenure:** Employees' length of time working at a role or company

Is it a Wave or a Particle?



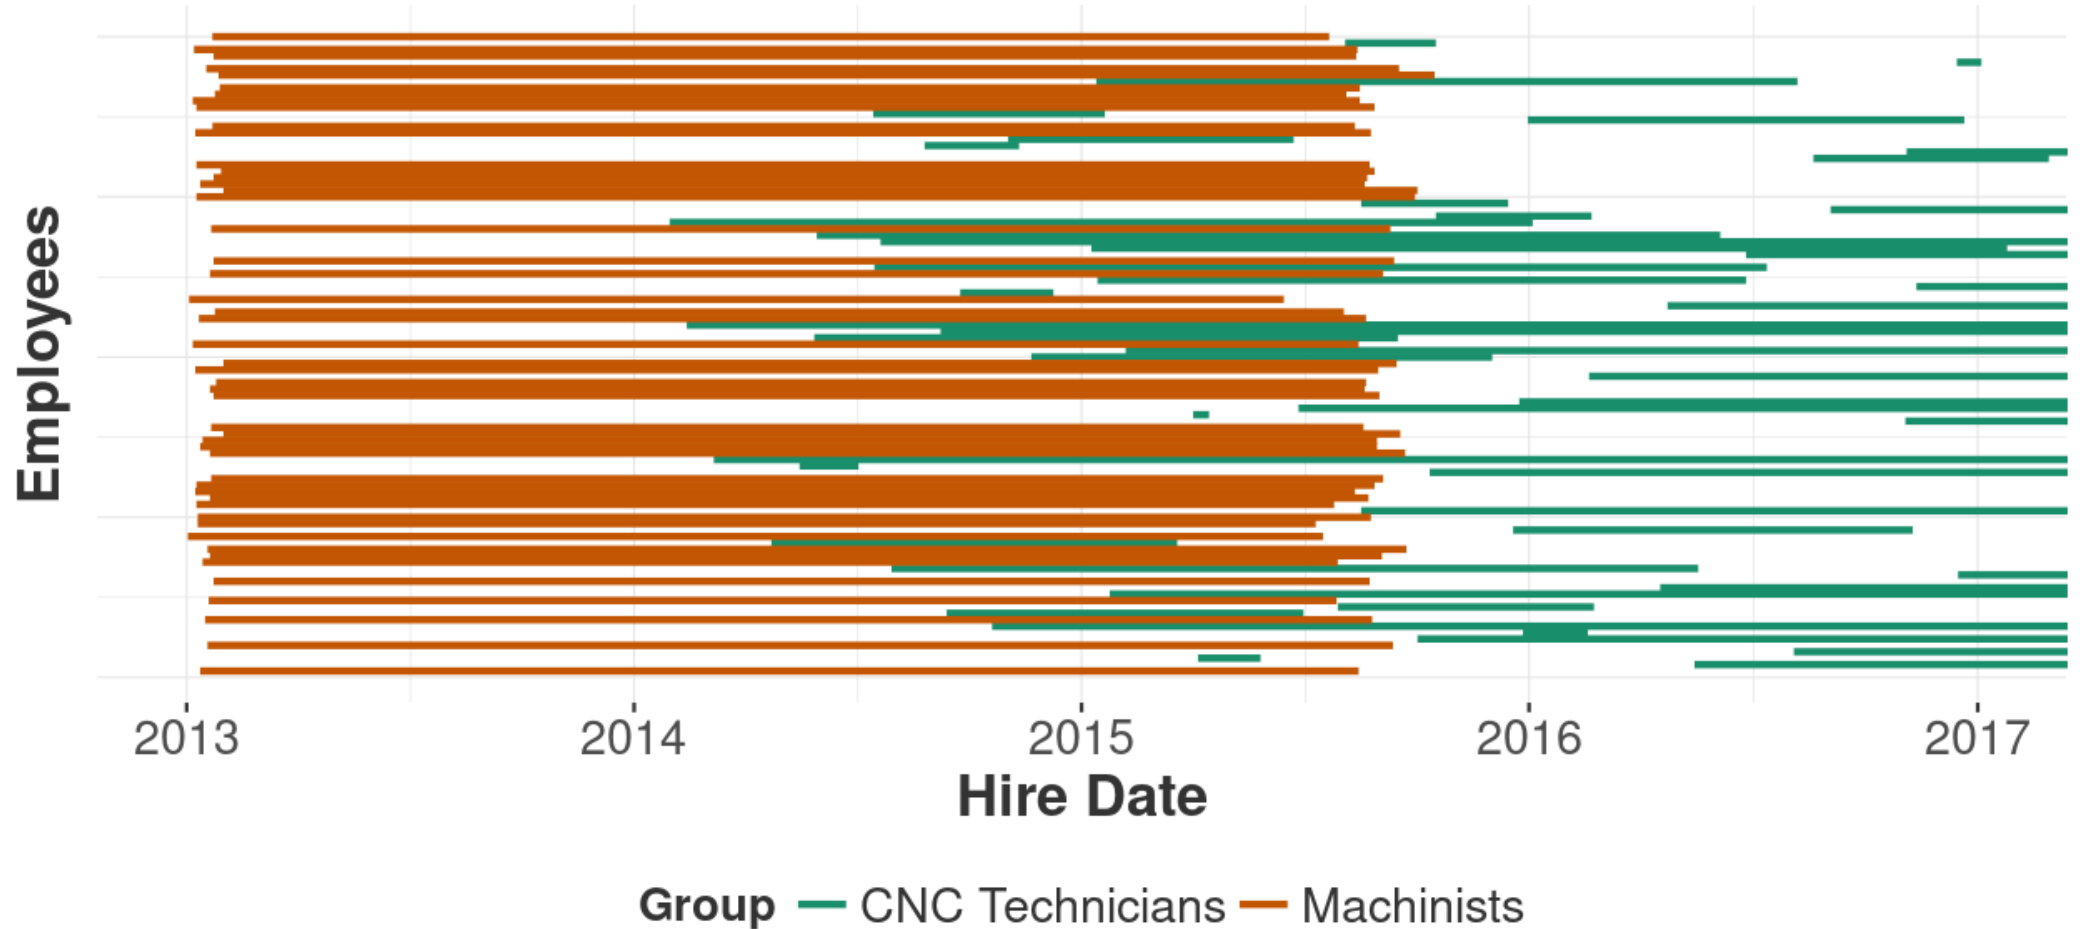
Low Tenure, High Turnover



High Tenure, Low Turnover



High Tenure, High Turnover (in 2015)




Low Tenure, Low Turnover (in Nov, Dec)



Main Problem: We Can't See the Future

- ▶ Everyone will terminate.
- ▶ Someday, somehow... But when?
- ▶ Technical Lingo: **Right Censoring**



Each Side Has Missing Information

Turnover Has No Information About Tenure

- ▶ Treats a temp the **same** as a seasoned veteran
- ▶ Ignores patterns - such as people quitting right after a bonus

Tenure Has No Information About Turnover


- ▶ High Tenure = staying power or retirement risk?

Each Side Varies Depending on When You Look

Annual Attrition Timing is Arbitrary

What if you could see attrition at every point in time?

- ▶ “One Year” only matters to accountants and astronomers
- ▶ Cumulative Breakeven matters more to the business



Survival 101

Survival Analytics

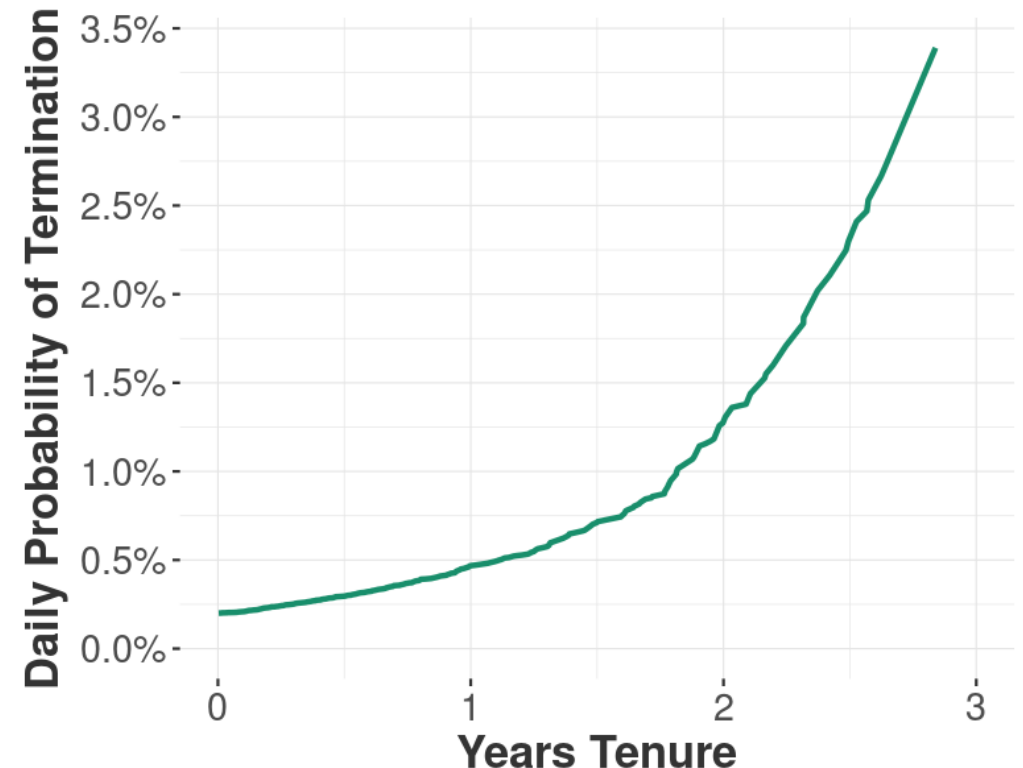
- ▶ **Medical Roots** - How long will a patient survive a disease?
- ▶ **Engineering Roots** - How long until a machine fails?
- ▶ **Social Science** - How long do people live?
- ▶ **Employment** - How long until an employee terminates? or promotes?

- ▶ **In General** - “Expected duration of time until one or more events happen”
- ▶ Artfully combines turnover and tenure

LOTS OF COMPETING JARGON AND SYNTAX FROM MANY APPLICATION DOMAINS:
FAILURE RATE, HAZARD RATE, FORCE OF MORTALITY, ...


Hazard Rate: Conditional Daily Risk of Termination

- ▶ Every day in tenure there is a (small) probability that the employee will terminate
- ▶ **Conditional on survival to the prior day**
- ▶ Like an actuarial life table
- ▶ Can rise at different times:
e.g. training, review or bonus time



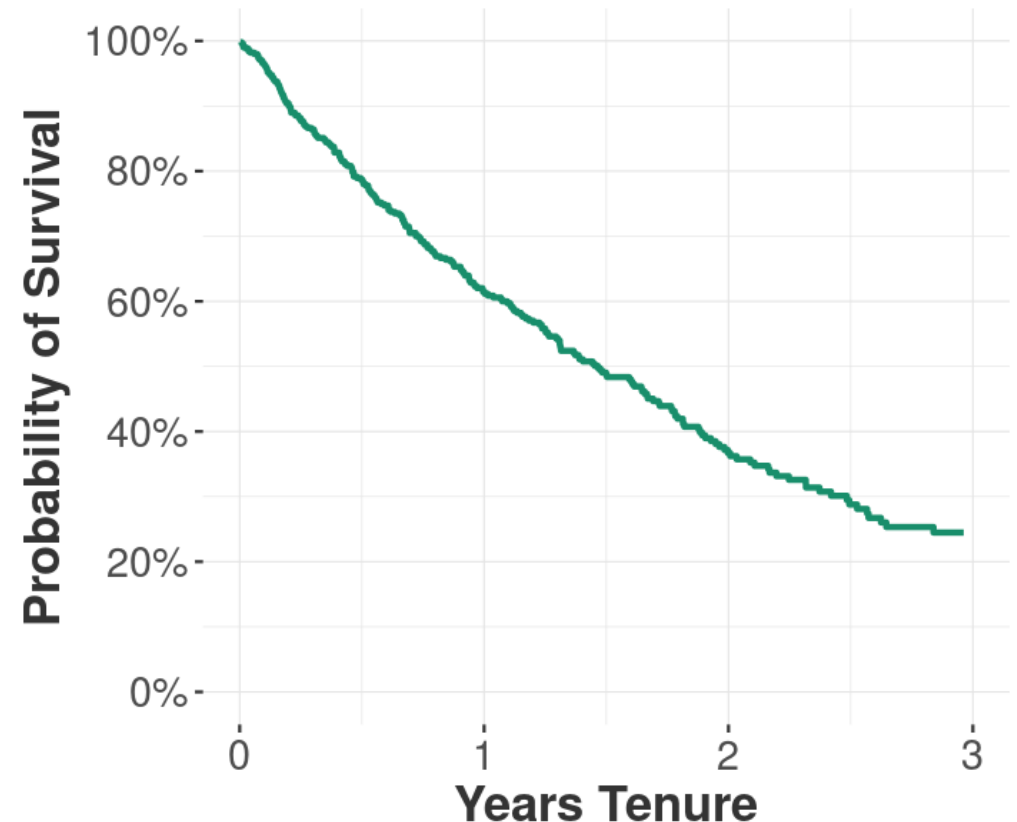
Survival Rate: Daily Probability of Being Employed at Time t

- ▶ How likely are you to be working here, t days after being hired?
- ▶ Near 100% on Day 1
(some people actually never show up)
- ▶ Downhill from there - usually not linear
- ▶ One role at a time




Survival Rate: Daily Probability of Being Employed at Time t

- ▶ How likely are you to be working here, t days after being hired?
- ▶ Near 100% on Day 1
(some people actually never show up)
- ▶ Downhill from there - usually not linear
- ▶ One role at a time
- ▶ *This* is the full picture
- ▶ Simple transformation of the Hazard Curve
 - ▷ Best to use statistical tool (R, SAS)
 - ▷ Apparently possible in Excel

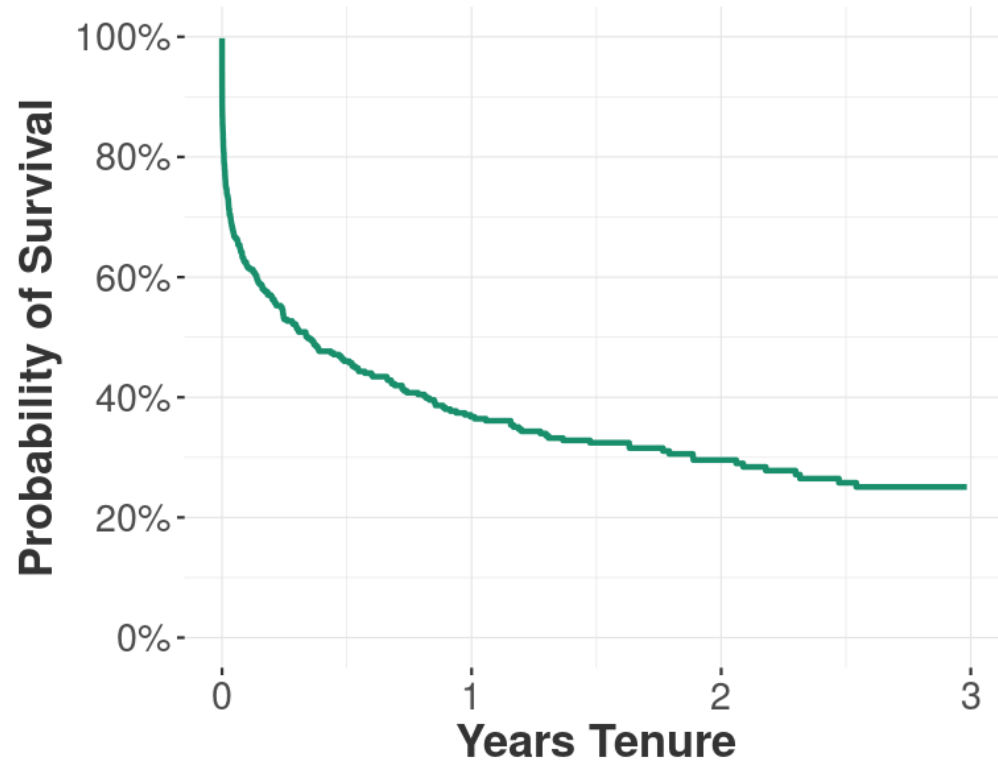


Attrition is Built Into Survival Curves

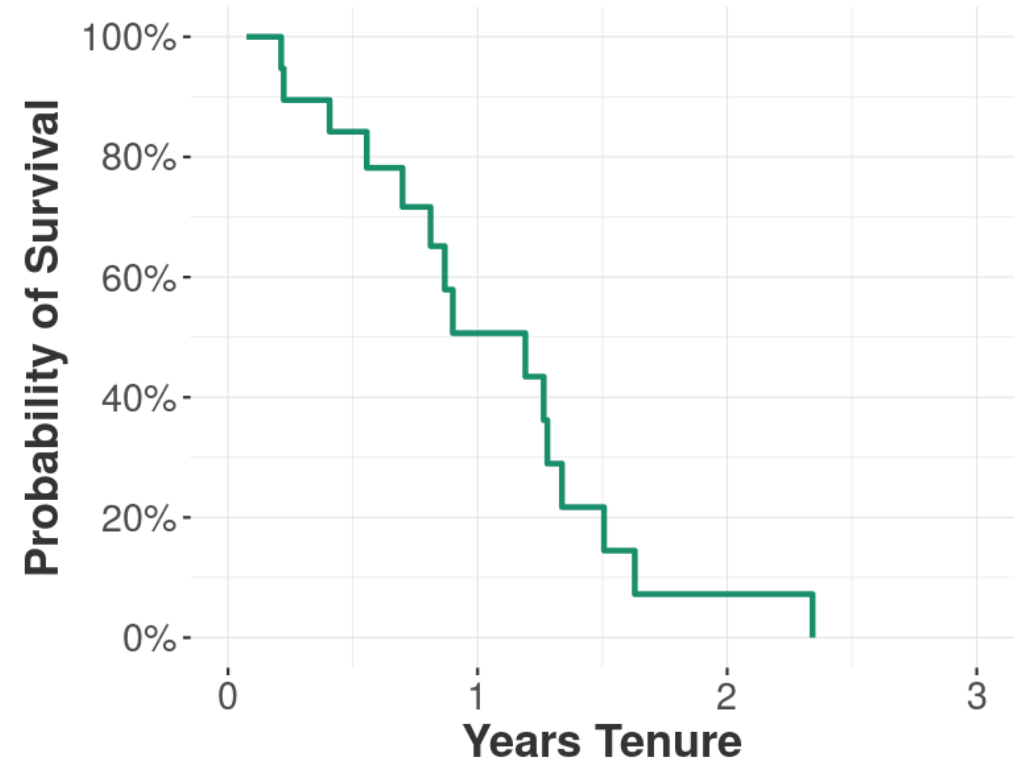


Many Shapes

Shelf Pattern




Sparse Pattern




Cliffs or Kinks in the Survival Curve

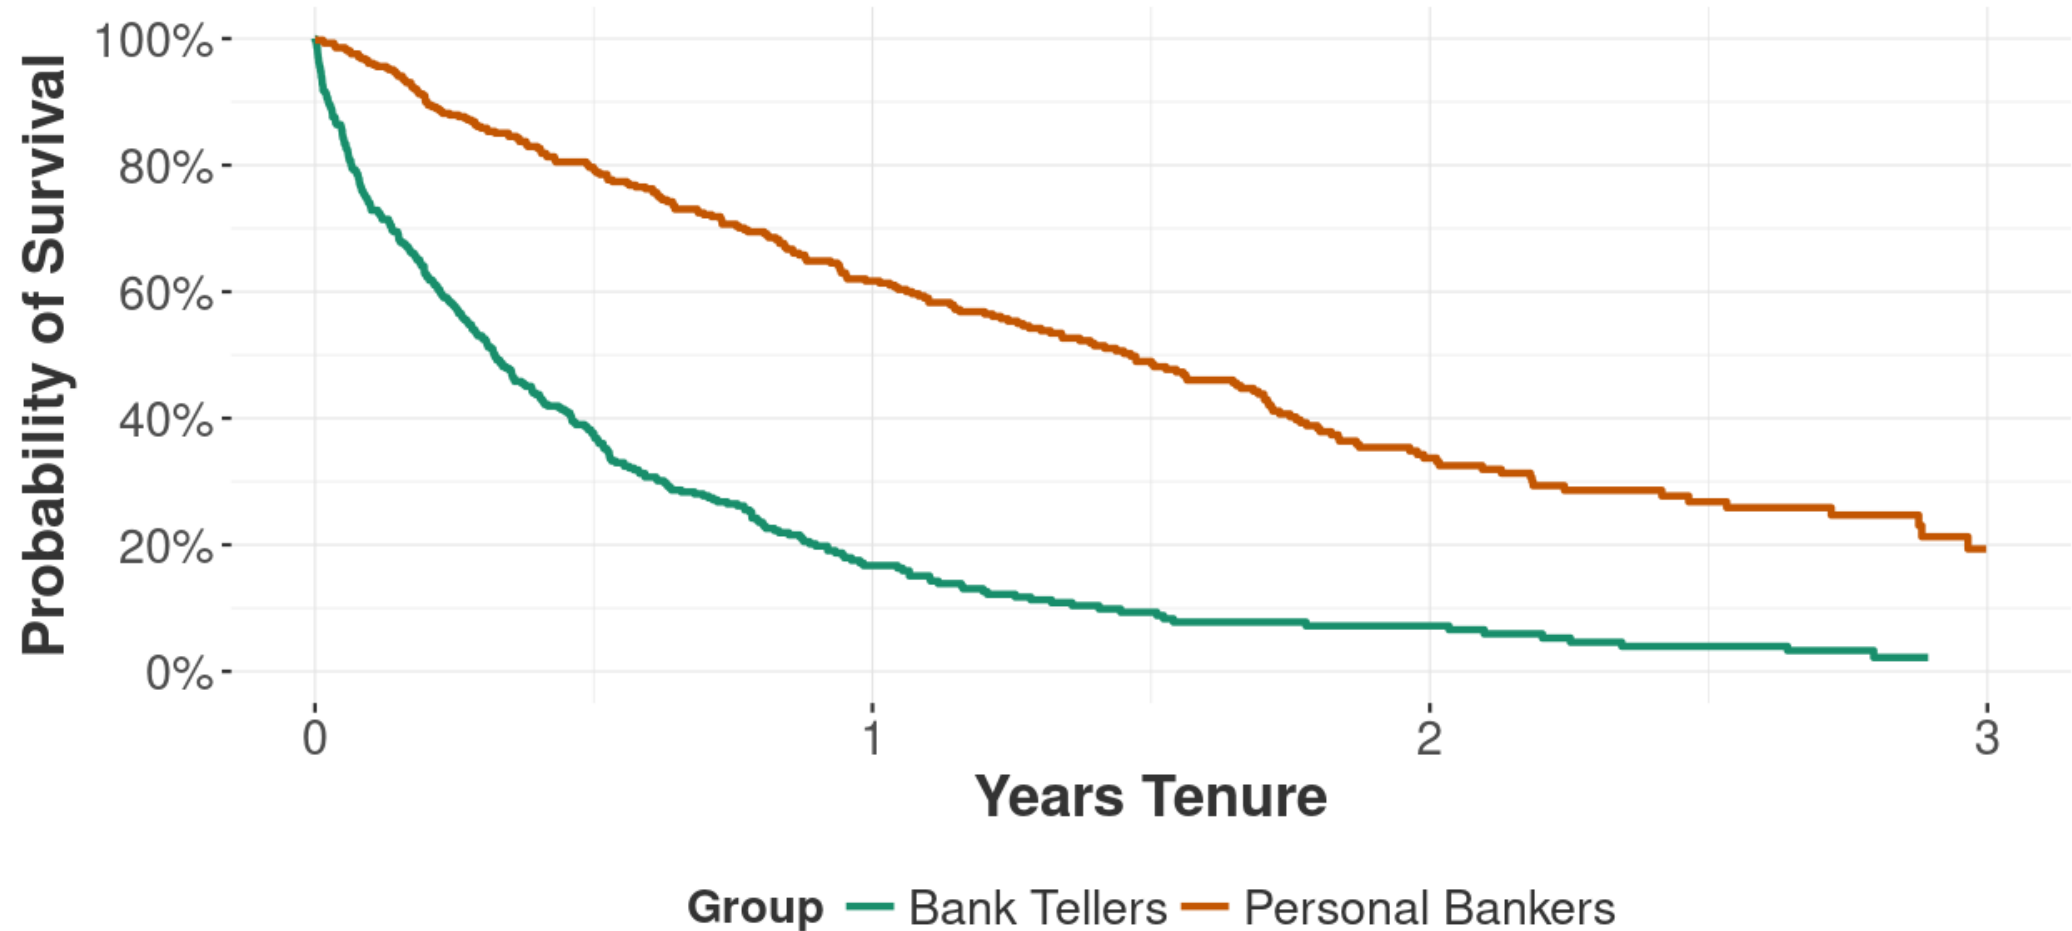
Heavy Terms During Training



Heavy Terms After 1-Year Bonus



Everyone Has a Survival Curve



Predicting Survival To the Individual Employee

Don't Use Logistic Methods to Predict Attrition!

Commonly Done:

- ▶ Predict one-year attrition with logistic
- ▶ Use tenure as a variable along with others

Common Result:

- ▶ “The biggest cause of termination is tenure”
 - ▷ Not an actionable coefficient
- ▶ Ignores nuances available to survival methods
- ▶ Mishandles current employee tenure
- ▶ Less accurate



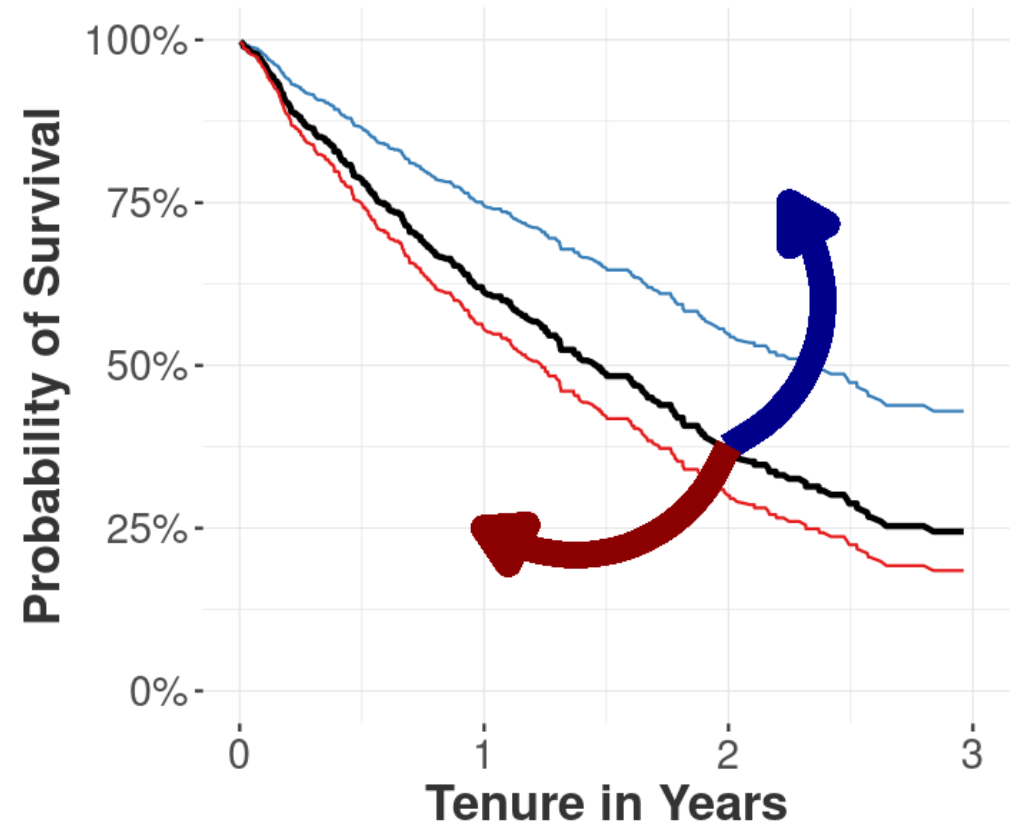
Proportional Hazard

Most Survival models are "Proportional Hazard"

- ▶ Start with Base Survival rate
- ▶ Convert to Base Cumulative Hazard rate
- ▶ Multiply Cumulative Hazard by a single factor (hence "Proportional Hazard")
- ▶ Linear changes to Hazard lead to a *rotation* of Survival Curve

Predictive Goal:

Predict the amount of Survival Curve rotation for each candidate

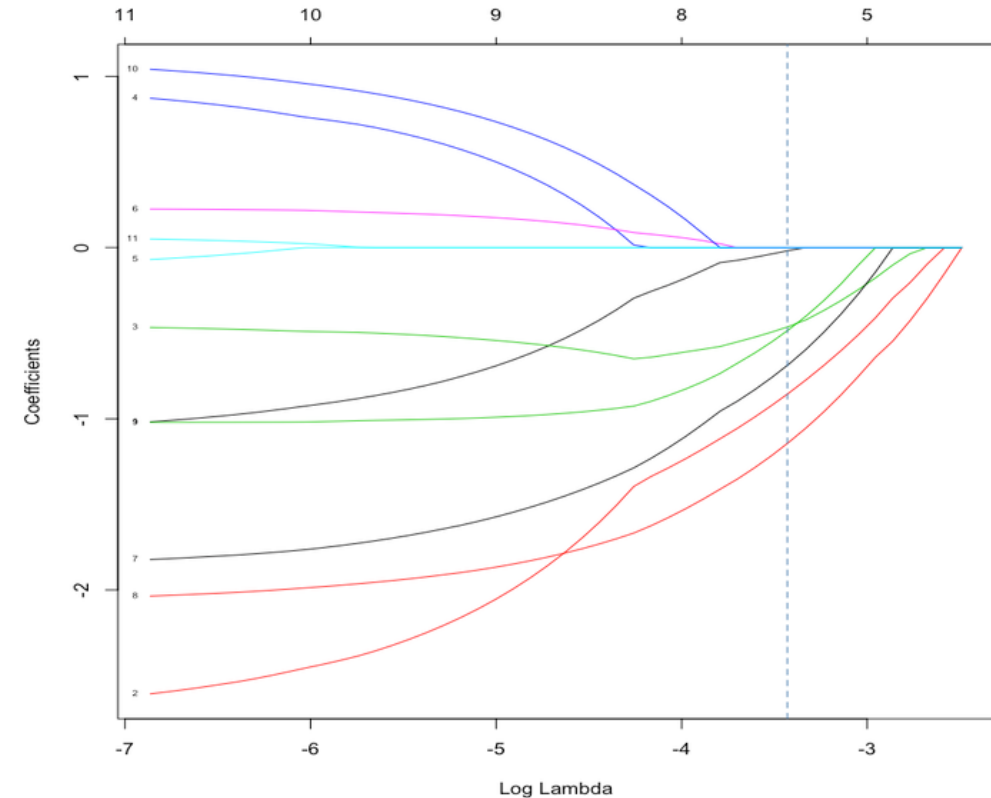


Proportional Hazard Prediction Methods

Predicting a Single Number

- ▶ Cox regression
 - ▷ Vanilla Cox
 - ▷ Stepwise
 - ▷ ElasticNet or LASSO
 - ▷ Variable selection with Random Forest
- ▶ Decision Trees
- ▶ Random Forests
- ▶ Neural Networks
- ▶ SVM, etc

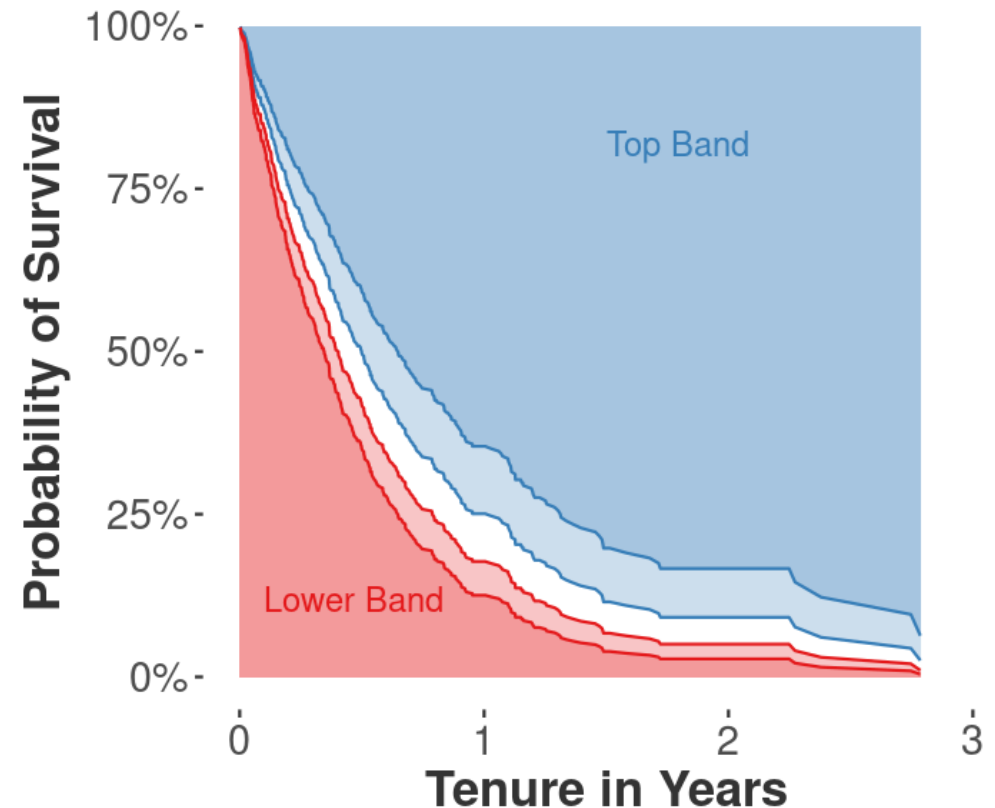
There are also other survival algorithms, including parametric methods



Build/Validate a Survival Model With Thresholds

The Model Predicts:

- ▶ “Dark Blue” hires will survive the longest
- ▶ “Light Blue” hires will survive above-median
- ▶ “Grey” survival curve is current median survival.
- ▶ “Light Red” and “Dark Red” candidates are not likely to survive.





Deeper Dive:



Predicting Survival with Proportional Hazard

Follow Demo Code on GitHub:

https://github.com/talentanalytics/class_survival_101/

Formulas and Jargon

Hazard Function: $h(t)$

- ▶ Probability of termination at time t , conditional on survival up to time t
- ▶ Ranges from 0 to 1 (typically very small)

Cumulative Hazard Function: $H(t) = \int_0^t h(x) dx = -\log(S(t))$

- ▶ Cumulative conditional probability of termination up to time t
- ▶ Ranges from 0 to 4ish (theoretically to Infinity)

Survival Function: $S(t) = e^{-H(t)}$

- ▶ Probability that termination will be later than time t
- ▶ Ranges from 1 to 0

Theoretically the formulas above are continuous measures, but in practice are discrete daily.

Calculating Survival in R

Four Simple Steps with R survival package:

1. Gather and Prepare the Data
2. Calculate Baseline Survival Function $S_0(t)$
3. Model Proportional Hazard with Cox Regression
4. Validate Model

Calculating Survival in R

Four Simple Steps with R survival package:

1. Gather and Prepare the Data
2. Calculate Baseline Survival Function $S_0(t)$
3. Model Proportional Hazard with Cox Regression
4. Validate Model
5. Deploy Model into hiring
6. Hire people likely to stay on the job longer
7. Save and earn money
 - ▶ Avoid attributing savings to soft factors like ripple effect, morale, engagement
 - ▶ Hard savings from reduced Replacement Cost
 - ▶ Hard earnings from increased Employee Lifetime Value

1) Gather Employee Data

```
> attr.data
```

	emp.id	hire.date	term.date	factor.x	factor.y	factor.z
	<int>	<dtm>	<dtm>	<dbl>	<dbl>	<dbl>
1	52	2016-02-23	<NA>	0.09555369	1.515028	102.3811
2	261	2014-10-30	<NA>	2.40144024	9.065855	100.5727
3	193	2013-12-25	2015-03-28	1.88463382	4.331642	105.6887
4	150	2013-03-27	2013-09-19	3.81625316	7.363454	110.6842
5	31	2015-03-25	<NA>	5.17392728	11.824601	110.0590

Simple Requirements

- ▶ Just hire.date and term.date is all you need from HRIS
 - ▷ Leave term.date as NA if the employee haven't terminated yet
- ▶ Merge with predictive input (independent) variables
 - ▷ e.g. assessment results, experience, CV detail, social media, semantic
 - ▷ Here we use factor.x, factor.y and factor.z
 - ▷ Are these variables available pre-hire?

Calculate Fields for Survival

```
> attr.data %>% dplyr::select(-dplyr::contains("factor"))
```

	emp.id	hire.date	term.date	is.term	end.date	tenure.years
	<int>	<dtm>	<dtm>	<lgl>	<dtm>	<dbl>
1	52	2016-02-23	<NA>	FALSE	2017-05-01	1.1854894
2	261	2014-10-30	<NA>	FALSE	2017-05-01	2.5023956
3	193	2013-12-25	2015-03-28	TRUE	2015-03-28	1.2539357
4	150	2013-03-27	2013-09-19	TRUE	2013-09-19	0.4818617
5	31	2015-03-25	<NA>	FALSE	2017-05-01	2.1026694

Simple Derivations

- ▶ `is.term`: Have they terminated? Is `term.date` NA?
- ▶ `end.date`: Either `term.date`, or censor date, depending on `is.term`
 - ▷ Could put `transfer.date` here with no `is.term`
 - ▷ *The last date we know the employee was at work*
- ▶ `tenure.years`: What is employee tenure from `hire.date` to `end.date`?
 - ▷ *All we know about the employee's tenure, as of censor date*

Dataset is Ready

```
> glimpse(attr.data)

Observations: 400
Variables: 14
$ label      <fctr> a, a, a, a, a, a, a, a, a, a, a, a, a, a, a, a, a, a, ...
$ hire.date  <dtm> 2013-04-20, 2016-01-14, 2013-09-20, 2014-05-17, 2016-...
$ term.date  <dtm> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, 2016-...
$ is.term    <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, ...
$ end.date   <dtm> 2017-05-01, 2017-05-01, 2017-05-01, 2017-05-01, 2017-...
$ tenure.years <dbl> 4.02975576, 1.29469635, 3.61029191, 2.95626651, 1.1463...
$ factor.x   <dbl> 7.44846858, 1.08376058, 2.81983987, 0.56880027, 0.7868...
$ factor.y   <dbl> 16.838313, 9.335225, 9.519951, 12.697147, 12.912898, 1...
$ factor.z   <dbl> 103.10732, 100.37171, 99.76382, 95.88645, 97.61509, 10...
$ emp.id     <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, ...
$ scale.x    <dbl> 2.41351313, -0.24754913, 0.47829958, -0.46285224, -0.3...
$ scale.y    <dbl> 1.31773198, -0.57718839, -0.53053544, 0.27187182, 0.32...
$ scale.z    <dbl> 0.3407242, -0.1849069, -0.3017096, -1.0467221, -0.7145...
$ is.training <lgl> TRUE, TRUE, TRUE, FALSE, FALSE, FALSE, TRUE, TRUE, TRU...
```

- ▶ Example is 400-row simulated dataset
- ▶ Scaled the input factors to consistent z-scale
- ▶ Separated into 80% training and 20% validation datasets

2) Calculate Baseline Survival ($S_0(t)$)

`survival::Surv` is a single object to hold Tenure and Turnover at once

```
> surv.obj <- survival::Surv(training.data$tenure.years, training.data$is.term)
```

- Needs “time variable” `tenure.years` and “event variable” `is.term`

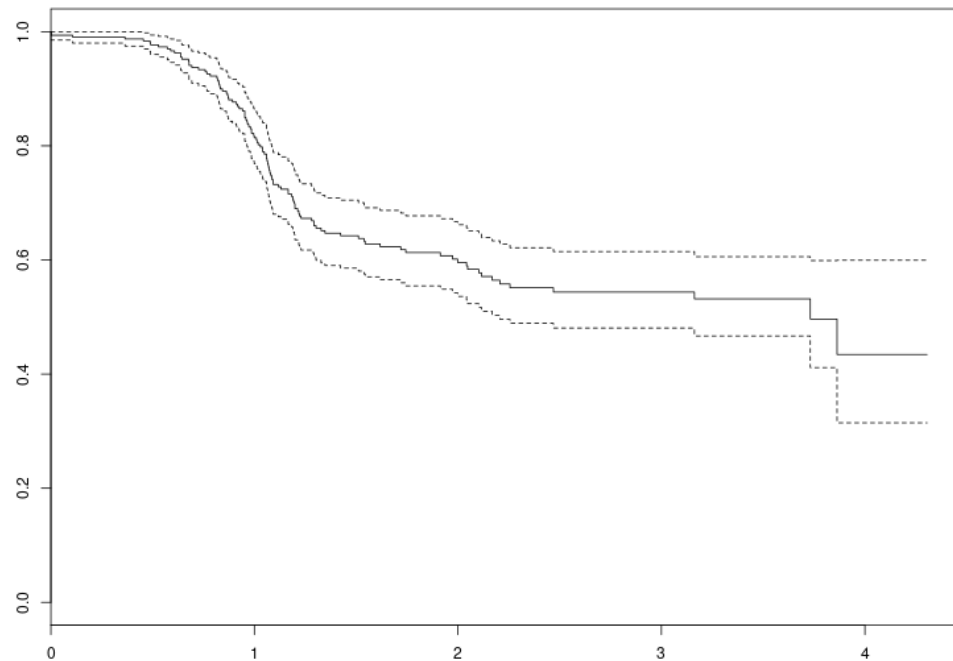
`survival::survfit` is the Kaplan-Meier Estimator of $S_0(t)$

```
> surv.fit <- survival::survfit(surv.obj ~ 1)
> summary(surv.fit)
Call: survfit(formula = surv.obj ~ 1)

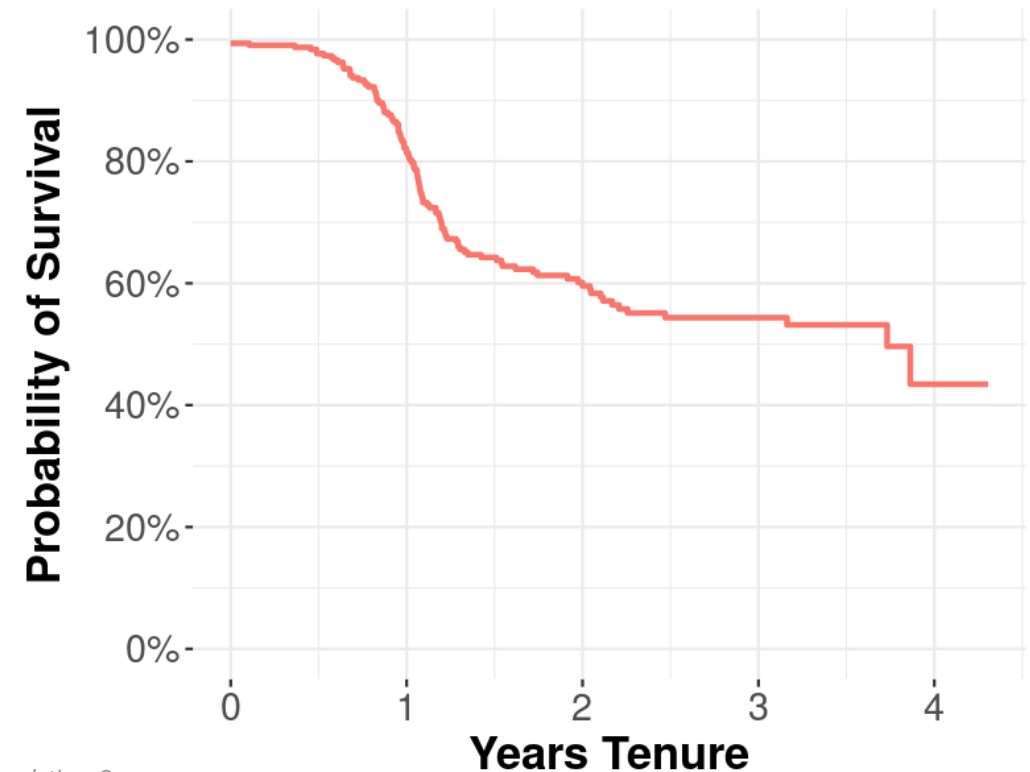
   time  n.risk  n.event  survival std.err lower 95% CI upper 95% CI
0.000    320      1    0.997 0.00312    0.991    1.000
0.249    308      1    0.994 0.00448    0.985    1.000
0.298    303      1    0.990 0.00554    0.980    1.000
0.320    298      1    0.987 0.00644    0.974    1.000
```

Plot Baseline Survival

Native Survival Plot



Crafted ggplot2 Plot



3) Back to Proportional Hazard

Consider a multiplier x that centers on 1


- ▶ Multiply $H_x(t) = H_0(t)x$
for proportional (linear) changes to Hazard
- ▶ Survival Curve **rotates** in response:
 $S_x(t) = e^{-H_0(t)x}$

Beware - inverse relationship:

- ▶ $x \leftarrow 1.1$ will increase $H_x(t)$, decrease $S_x(t)$
- ▶ $x \leftarrow 0.9$ will decrease $H_x(t)$, increase $S_x(t)$

So we just need to predict x to predict Survival

- ▶ One number drives the entire curve, based on baseline $H_0(t)$



Model Proportional Hazard with Cox Regression

```
> cox.model <- survival::coxph(formula = surv.obj ~ scale.x + scale.y + scale.z,  
                               data = training.data)
```

```
> cox.model
```

Call:

```
survival::coxph(formula = surv.obj ~ scale.x + scale.y + scale.z,  
                data = training.data)
```

	coef	exp(coef)	se(coef)	z	p
scale.x	-0.4298	0.6506	0.0918	-4.68	0.00000028
scale.y	-0.3204	0.7258	0.0957	-3.35	0.00082
scale.z	-0.1967	0.8215	0.0984	-2.00	0.04558

```
Likelihood ratio test=38.2 on 3 df, p=0.000000025  
n= 320, number of events= 125
```

4) Predict validation.data with New Cox Model

```
> cox.pred <- predict(cox.model, newdata = validation.data, type = "lp")
```

- ▶ **Note:** we built model with training.data, now we predict against validation.data
 - ▷ Our model has never seen anything in validation.data
 - ▷ Great test to see how closely our predictions match known outcomes
- ▶ The predict.coxph() function returns 5 different types of predictions
 - ▷ We want type = "lp" which is log(multiplier) to baseline $H_0(t)$
 - ▷ Other types won't work with survivalROC()

How does ROC Work with Survival?

```
> roc.obj <- survivalROC::survivalROC(Stime = validation.data$tenure.years,  
                                     status = validation.data$is.term,  
                                     marker = cox.pred,  
                                     predict.time = 1,  
                                     lambda = 0.003)
```

```
> roc.obj$AUC  
[1] 0.7102442
```

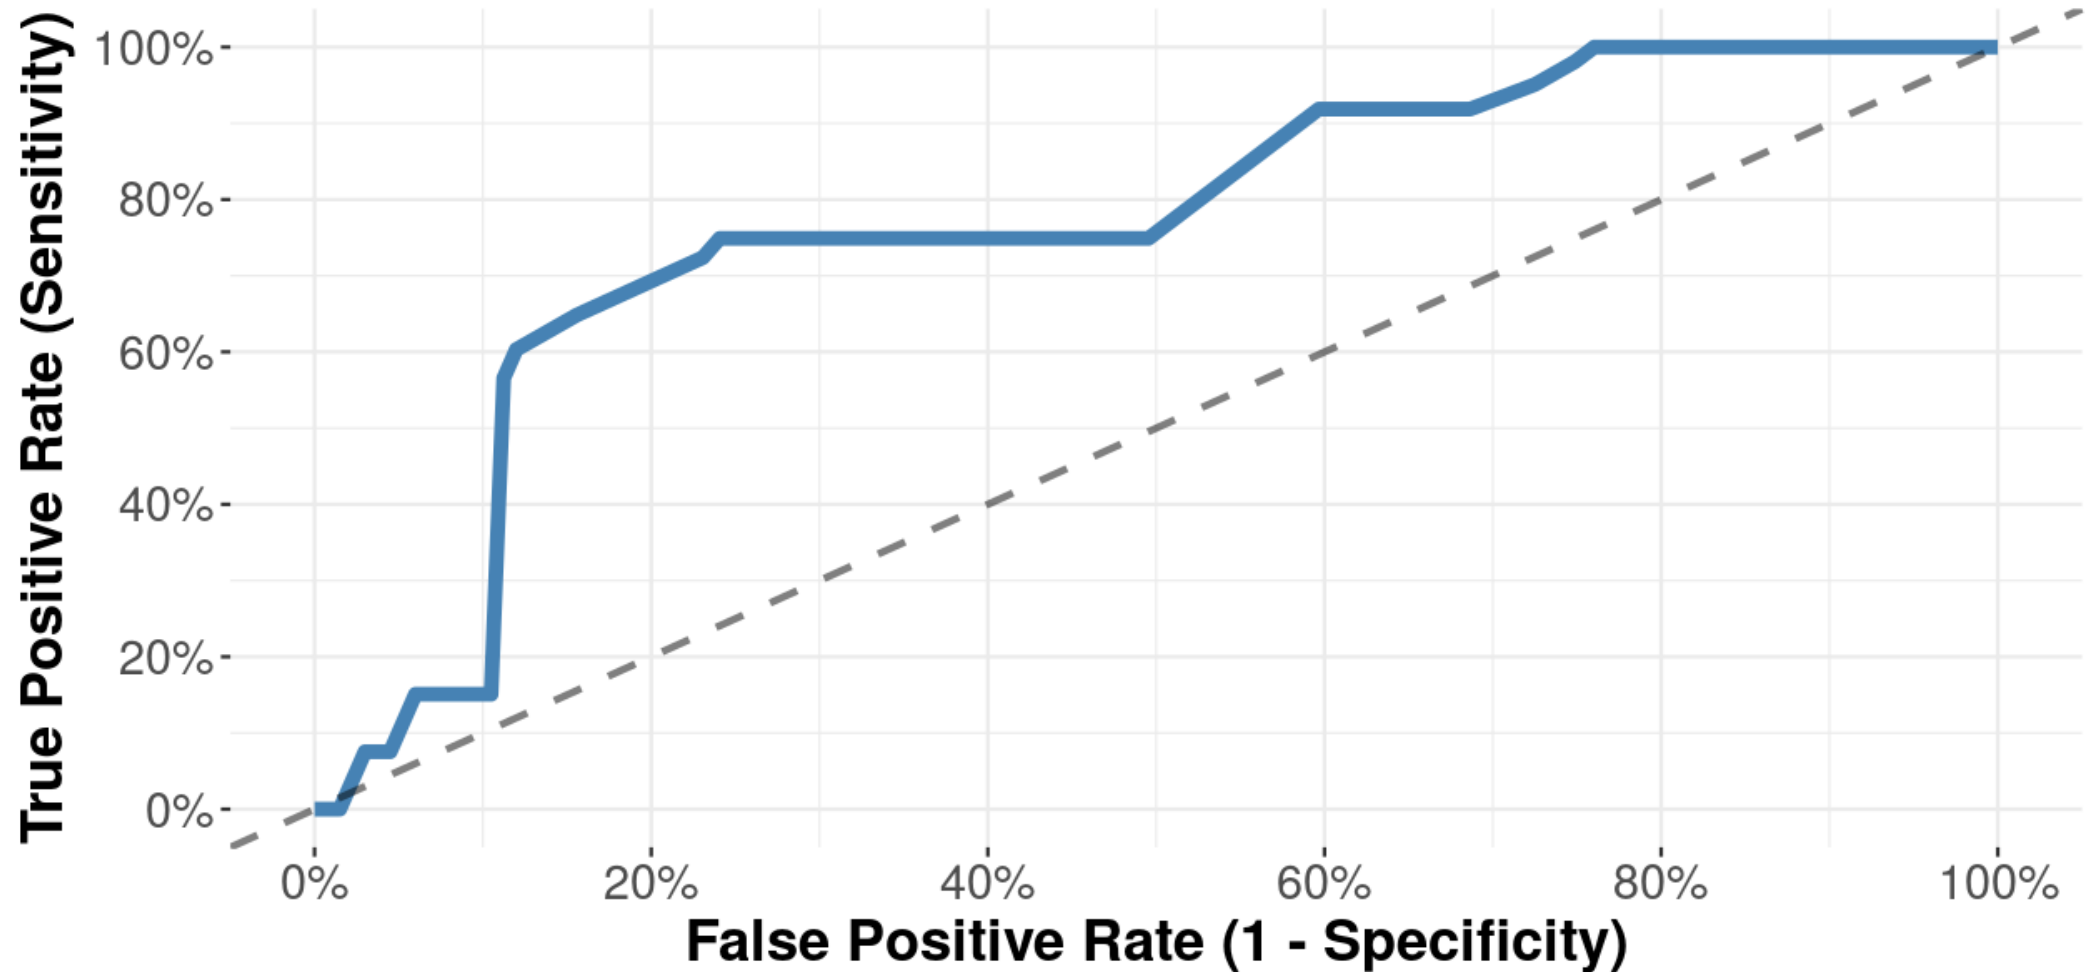
ROC typically evaluates classification methods

- ▶ So, we make this a classification problem
- ▶ **Classify:** Was the employee terminated at $t = 1$ or not?
- ▶ Compare actual vs. predicted classifications at all cutpoints

Our AUC is 0.71

- ▶ Not bad for a censored attrition prediction!

Plot AUC with ggplot2



Don't Get Fooled by Randomness

We are predicting AUC, too

- ▶ Let's try predicting 1000 random samples of the same underlying pattern

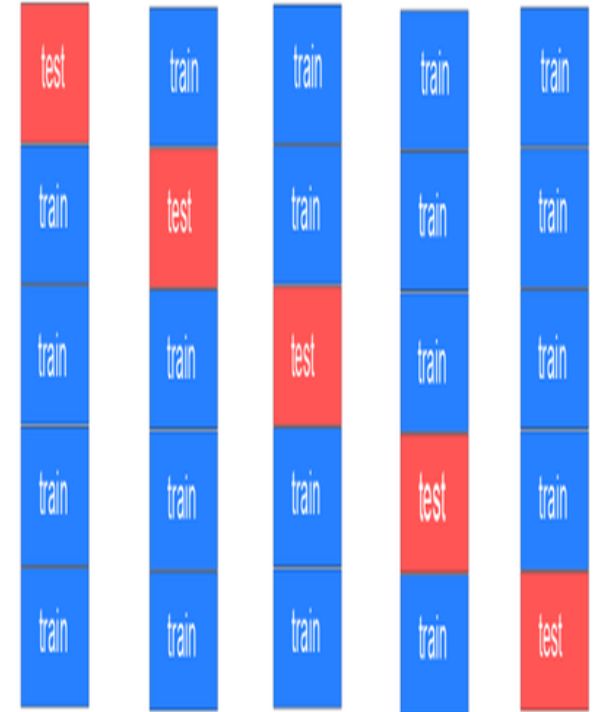
```
> test.repl <- replicate(1000, demoPrediction(verbose = FALSE))  
  
> summary(test.repl)  
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   
 0.3470  0.6191  0.6833  0.6771  0.7356  0.9446
```

- ▶ AUC generally fell between 0.619 and 0.735, but 50% were outside this range
- ▶ The likely AUC is closer to 0.68 - still good
- ▶ What if you got a **lucky** (or unlucky) 80/20 validation split?

Multiple Cross-Validation

Raises the Bar Against Getting Fooled

- ▶ Typically run 20-100 CV iterations
- ▶ Can consume lots of core-hours if nested for hyperparameters



Download and Experiment With this Code

- ▶ Download demo code and doc from https://github.com/talentanalytics/class_survival_101/
- ▶ Clone it, fork it, run it
- ▶ Send bug fixes
- ▶ Let me know how it goes



Questions and Discussion

Pasha Roberts

pasha@talentanalytics.com

Follow Demo Code on GitHub:

https://github.com/talentanalytics/class_survival_101/